

ISOM 3360

ISOM 3360 - Data Mining for Business Analytics

Spring 2020, HKUST

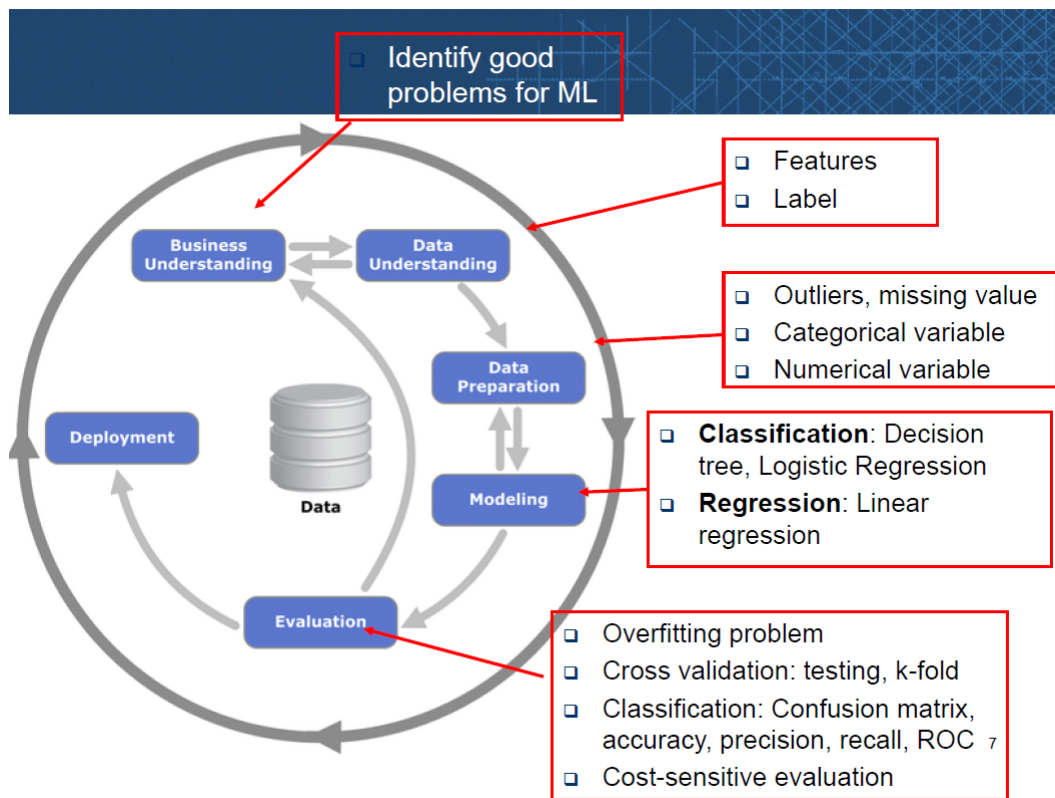
This repository contains code of labs and assignments.

For reference only, please DO NOT COPY.

ISOM 3360

- Overview of Data Mining
 - Strengths and weakness
 - Data preparation
 - Missing values handling
 - Feature transformation
 - Balancing
 - Modeling
 - Train-test split
 - Cross validation
 - Overfitting
 - Decision threshold
 - Model Evaluation
 - Confusion Matrix
 - Precision-Recall Curve
 - ROC Curve
 - AUC
 - Cost-sensitive classification
 - Prediction bias
- Decision Tree
 - Metrics
 - Gini Impurity
 - Information Gain
 - ID3 Algorithm
 - Decision boundaries
 - Pros and Cons
- Linear Regression
 - Basics
 - Lasso Regression
 - Evaluation
 - Pros and Cons
- Logistic Regression
 - Basics
 - Decision boundaries

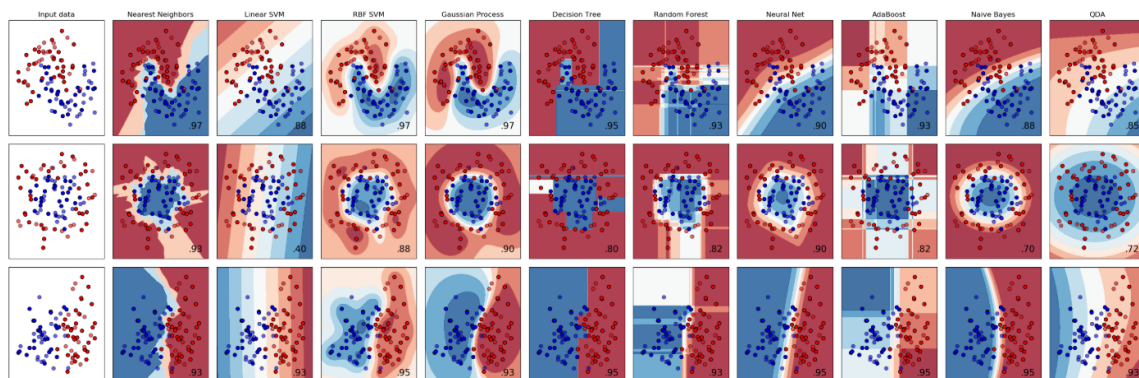
Overview of Data Mining



Supervised learning (prediction of labels)

Unsupervised learning (relationship mining)

Classification decision boundaries:



Strengths and weakness

ML tends to work well when

- Learning a simple concept
- Lots of (good quality) data available

ML tends to work poorly when

- Learning complex concepts from a small amount of data
- Performing on new types of data
 - e.g. photos taken in big hospitals and small hospitals are different types, because small hospitals may rotate.

Data preparation

- Size: $> 10 \times \text{number of features}$

- Quality: label errors, noisy features, unreliable source

Outliers: $> 3 \sigma$ away from the mean

Missing values handling

- Ignore instances/features
- Fill in:
 - Global constant (Null or 0)
 - **Feature mean**
 - Feature mean for samples in each class
 - Feature mode/medium/...

Feature transformation

- Categorical: **one-hot encoding**
- Numerical:
 - **normalization** (after visualization)
 - scaling to range: e.g. min-max: $(x - x_{min}) / (x_{max} - x_{min})$
 - clipping: capped above/below a threshold \Rightarrow bell shape
 - log scaling: $\log(x)$ for long-tailed data
 - z-score scaling: $(x - \mu) / \sigma$ for normal data
 - **discretization**
 - binning

Balancing

Degree of Imbalance	Proportion of the Minority Class
Mild	20% - 40%
Moderate	1% - 20%
Extreme	$< 1\%$

- Resample
- Generate synthetic data for the minority class

Modeling

Train-test split


Avoid data leakage: never use testing data in data preparation (e.g. `fillna`) or modeling

Testing set:

- Large enough to yield statistically meaningful results
- Is representative of the data set as a whole

Cross validation

k-fold ($k = 10$)



Overfitting

Noise or random fluctuations in the training data is picked up and learned.

Decision threshold

Probabilistic predictions: e.g. $p \geq 0.5 \Rightarrow y = 1$

Determine the confusion matrix and ROC curve.

Model Evaluation

Benchmark: the naïve learner e.g. random guess, **majority class classifier**

Confusion Matrix

		Predicted class	
		+	-
Actual class	+	True + (TP)	False - (FN)
	-	False + (FP)	True - (TN)

Accuracy = $(TP + TN) / (TP + TN + FP + FN)$

Error Rate = $1 - \text{Accuracy}$

Precision = $TP / (TP + FP)$


Recall (TPR) = $TP / (TP + FN)$

Fall-out (FPR) = $FP / (FP + TN)$

A confusion matrix is determined by its **decision threshold**.

Threshold \uparrow , Precision \uparrow , Recall \downarrow

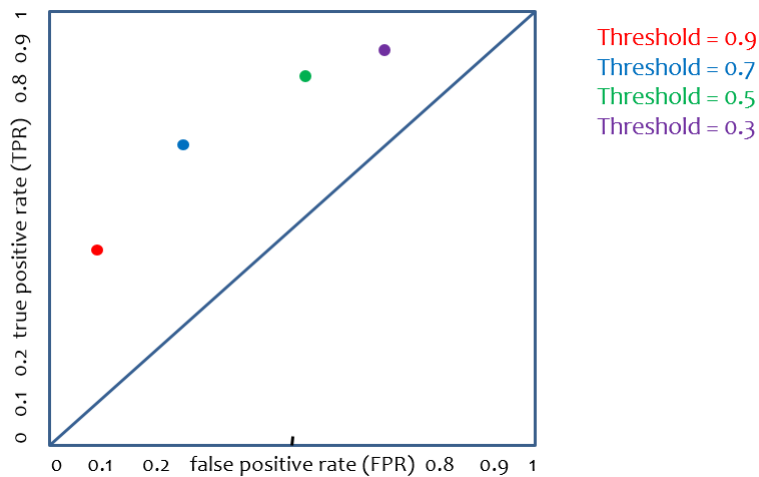
Precision-Recall Curve



high precision → few results

Useful for imbalanced datasets where the positive class is the minority (e.g. web search).

ROC Curve



- Diagonal: random guess, AUC = 0.5
 - (1, 1): $y = 1$
 - (0, 0): $y = 0$
- (0, 1): perfect model, AUC = 1
- (1, 0): perfect model = Not(y)

AUC

Area Under (the ROC) Curve: between 0.5 (random) and 1 (perfect).

- aggregate measure of performance across all possible thresholds
- probability that the model ranks a random positive example higher than a random negative example

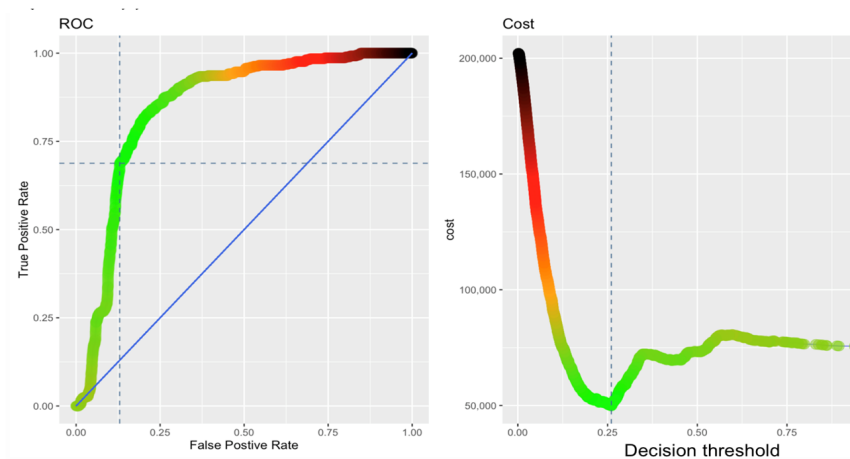
Cost-sensitive classification

$$\text{Cost} = \text{FP} \times (\text{cost of FP}) + \text{FN} \times (\text{cost of FN})$$

When the costs of FP (Type I error) and FN (Type II error) are different,

Minimize the cost → Minimize the type of error with the higher cost.

Examples: credit card applications, spam filtering, medical diagnosis, customer churn



Prediction bias

Bias = average of predictions – average of actual labels

Significant non-zero bias \Rightarrow Error in model

Decision Tree

Find features with high predictability.

Structure:

- upside down if-else tree
- **root** node contains all training examples
- **leaf** node
- numerical features may be discretized

Metrics

Gini Impurity

Information Gain

Entropy: uncertainty/impurity in the dataset; always non-negative

$$H(S) = \sum_{x \in X} -p(x) \log_2 p(x)$$

where $p(x)$ is the proportion of class x in the dataset S .

Information Gain: decrease in entropy after a split

$$IG(A, S) = H(S) - \sum_{t \in T} p(t) H(t)$$

where T are the subsets created from splitting S by feature A .

ID3 Algorithm

Only for classification, only handles categorical features.

1. Calculate the entropy of every feature using the dataset S .
2. Split the set S into subsets using the feature for which the information gain is maximum.

3. Make a decision tree node containing that feature, divide the dataset by its branches, and calculate the entropy of each branch
4.
 - If a branch has entropy of 0, it's a leaf node.
 - If `depth = max_depth` or `size < min_samples_split`, it's a leaf node.
 - Else:
 - If additional splits obtain no information gain, it's a leaf node.
 - Else, go back to Step 2.

Decision boundaries

Parallel to the axes.

Pros and Cons

Pros

- easy to understand, implement, and use (good for non-DM-savvy stakeholders)
- computationally cheap
- require little data preparation (no need for one-hot encoding)

Cons

- can overfit, need pruning
- unstable structure (one minor change can make the tree completely different)
- decision boundaries are parallel to the axes
- accuracy is mediocre

Linear Regression

Basics

$$\begin{aligned}
 h_{\mathbf{w}}(\mathbf{x}) &= \mathbf{w}^T \mathbf{x} \\
 &= w_0 + \sum_{i=1}^n w_i x_i \\
 \mathbf{w} &= \underset{\mathbf{w}}{\operatorname{argmin}} \sum (h_{\mathbf{w}}(x) - y)^2 \\
 &= (X^T X)^{-1} X^T y
 \end{aligned}$$

For univariate linear regression,

$$\begin{aligned}
 w_1 &= \frac{\sum (x - \bar{x})(y - \bar{y})}{\sum (x - \bar{x})^2} \\
 w_0 &= \bar{y} - w_1 \bar{x}
 \end{aligned}$$

Lasso Regression

$$\mathbf{w} = \underset{\mathbf{w}}{\operatorname{argmin}} \sum (h_{\mathbf{w}}(x) - y)^2 + \lambda \sum |w|$$

L1 Regularization: $\lambda \uparrow$, fitness \uparrow , model complexity \uparrow .

- Penalize models with too many features \Rightarrow Generate **sparse models**
- Penalize features with large coefficients

Evaluation

Root Mean Squared Error (RMSE): penalize large prediction errors

$$\text{RMSE} = \sqrt{\frac{1}{N} \sum (\hat{y} - y)^2}$$

Mean Absolute Error (MAE):

$$\text{MAE} = \frac{1}{N} \sum |\hat{y} - y|$$

Both RMSE and MAE are sensitive to outliers.

Pros and Cons

Pros

- Simple and easy to interpret

Cons

- Oversimplifies problems by assuming linearity
- Sensitive to outliers

Logistic Regression

Basics

Softmax function

$$\sigma = \frac{\exp(x_k)}{\sum \exp(x)}$$

Conditional probabilities

$$\begin{aligned} P(y = k \mid \mathbf{x}) &= \sigma(\mathbf{w}_k^T \mathbf{x}) \\ &= \frac{\exp(\mathbf{w}_k^T \mathbf{x})}{\sum \exp(\mathbf{w}^T \mathbf{x})} \end{aligned}$$

For binary classification, we have

$$\begin{aligned} P(y = 1 \mid \mathbf{x}) &= \sigma(\mathbf{w}_1^T \mathbf{x}) \\ &= \frac{\exp(\mathbf{w}_1^T \mathbf{x})}{\exp(\mathbf{w}_0^T \mathbf{x}) + \exp(\mathbf{w}_1^T \mathbf{x})} \\ &= \frac{1}{1 + \exp[(\mathbf{w}_1 - \mathbf{w}_0)^T \mathbf{x}]} \\ &= \frac{1}{1 + \exp(\mathbf{w}^T \mathbf{x})} \end{aligned}$$

where $\mathbf{w} = \mathbf{w}_1 - \mathbf{w}_0$.

So, we have the sigmoid function

$$s = \frac{1}{1 + \exp(-x)}$$

And we can get the conditional probability from

$$P(y = 1 \mid \mathbf{x}) = s(\mathbf{w}^T \mathbf{x})$$

Decision boundaries

- Univariate: parallel to y-axis
- Multivariate: linear but not parallel to y-axis